

RoSteALS Implementation and Robustness

Eugene Francisco

June 2026

Overview

During a one week break between Stanford's Spring and Summer quarters, I reimplemented Bui et al.'s¹ RoSteALS technique for image watermarking.

As an overview of the technique, RoSteALS embeds watermarks into a cover image by using an image autoencoder and embedding information within the cover image's latent representation. More formally, let E, G be an autoencoder. For a high fidelity image $c \in \mathbb{R}^{H \times W \times C}$, let $E(c) = z \in \mathbb{R}^{H' \times W' \times C}$ be c 's latent representation. For a latent variable $z \in \mathbb{R}^{H' \times W' \times C}$, let $G(z) \in \mathbb{R}^{H \times W \times C}$ be the generated image from the auto encoder.

How does RoSteALS watermark images? Let $c \in \mathbb{R}^{H \times W \times C}$ be an image we intend to watermark and $z = E(c)$ be its encoding. Let $s \in \mathbb{R}^\ell$ be a secret message we intend to embed in c . RoSteALS learns two neural networks:

- (i) First, it learns a message encoder F_θ which outputs $\delta = F_\theta(s) \in \mathbb{R}^{H' \times W' \times C}$. The watermarked latent is then $z + \delta$ and the watermarked image is $\tilde{c} = G(z + \delta)$.
- (ii) Second, it learns a secret decoder D_φ which outputs $\tilde{s} = D_\varphi(\tilde{c})$, a reconstruction of the original message.

The neural networks are parametrized by learnable parameters θ and φ . Note that the autoencoder E, G is not learned and is fixed in advance. They are trained using two objectives:

- (a) $\mathcal{L}_{\text{quality}} = \mathcal{L}_{\text{LPIPS}}(\tilde{c}, c) + \alpha \|\gamma(\tilde{c}) - \gamma c\|_2^2$, where $\mathcal{L}_{\text{LPIPS}}$ is the LPIPS image reconstruction loss and γ is the mapping function from RGB to YUV.
- (b) $\mathcal{L}_{\text{recovery}} = \text{BCE}(s, \tilde{s})$ where BCE is the binary cross entropy between s and $\tilde{s}$.

The training loss is then

$$\mathcal{L} = \beta \mathcal{L}_{\text{quality}} + \mathcal{L}_{\text{recovery}}.$$

¹<https://arxiv.org/pdf/2304.03400>

Implementation Details

The purple text corresponds to implementation details from Bui et al. The olive text corresponds to my own specific implementation details. The VQ-GAN vq-f4 autoencoder is used; this uses a latent dimension of $H' = H/4$ and $W' = W/4$. I fixed hyperparameters $\alpha = 1.5$ and β is scaled dynamically (see below). I use the AdamW optimizer with learning rate $2 \cdot 10^{-5}$. I trained on the COCO training dataset from 2017 with 100,000 images of size 256×256 and a message length of $\ell = 50$. I used curriculum learning to train with the following training schedule:

- (a) *Warmup phase 1*: I fix $\beta = 0.1$ for the entirety of the phase and train F_θ and D_φ until the bit accuracy reaches 0.9. Training is done with a fixed subset of 8 randomly selected images from the total training set and a minibatch size of 4.
- (b) *Warmup phase 2*: I reveal an additional 50,000 images. β continues to be fixed at 0.1. Batch size is bumped to 8. This phase ends once bit accuracy reaches 0.8.²
- (c) *Exposure phase 1*: Over the course of 5,000³ training iterations, I linearly increase β from 0.1 to 10.
- (d) *Exposure phase 2*: I reveal the full 100,000 image training set to the models. This phase ends once bit accuracy reaches 0.98.
- (e) *Noising phase*: **TODO** blah blah

²In my experiments, the transition from *warmup phase 1* to *warmup phase 2* is typically followed by a drastic drop in bit accuracy, hence the lowered threshold.

³From my experiments, 5,000 corresponded roughly to a 10 percentage point increase in bit accuracy